

Performance Modelling - RISC-V processor

This project will require you to implement cycle-accurate simulators of a 32-bit RISC-V processor in C++ or Python. The skeleton code for the assignment is given in file (NYU_RV32I_6913.cpp or NYU_RV32I_6913.py).

The simulators should take in two files as inputs: imem.txt and dmem.txt files
The simulator should give out the following:

- cycle by cycle state of the register file (RFOutput.txt)
- Cycle by cycle microarchitectural state of the machine (StateResult.txt)
- Resulting dmem data after the execution of the program (DmemResult.txt)

The imem.txt file is used to initialize the instruction memory. The dmem.txt file is used to initialize the data memory of the processor. The instructions in the imem.txt file are 32-bit and the data in the dmem.txt file are 32-bit. The instructions in the imem.txt file are in “Big-Endian” format (the most significant byte is at the lowest address).

The instructions to be simulated are of the following types:

Bit	31			7, 6	0
R-type	funct7 (7 bits)				Opcode (7 bits)
I-type	imm (31 bits)				Opcode
S-type	imm[11:5]				Opcode
B-type	imm[12, 10:5]				Opcode
U-type					Opcode
J-type					Opcode

The simulator should support the following set of instructions.

Mnemonic	Type	Full Name	Pseudocode	Details
ADD	R	Addition	$rd = rs1 + rs2$	Store the result of $rs1 + rs2$ in register rd .
SUB	R	Subtraction	$rd = rs1 - rs2$	Store the result of $rs1 - rs2$ in register rd .
XOR	R	Bitwise XOR	$rd = rs1 \wedge rs2$	Store the result of $rs1 \wedge rs2$ in register rd .
OR	R	Bitwise OR	$rd = rs1 \mid rs2$	Store the result of $rs1 \mid rs2$ in register rd .

AND	R	Bitwise AND	$rd = rs1 \& rs2$	Store the result of $rs1 \& rs2$ in register rd .
ADDI	I	Add Immediate	$rd = rs1 + \text{sign_ext}(imm)$	Add the sign-extended immediate to register $rs1$ and store in rd . Overflow bits ignored.
XORI	I	XOR Immediate	$rd = rs1 \wedge \text{sign_ext}(imm)$	Bitwise XOR the sign-extended immediate to register $rs1$ and store result in rd .
ORI	I	OR Immediate	$rd = rs1 \mid \text{sign_ext}(imm)$	Bitwise OR the sign-extended immediate to register $rs1$ and store result in rd .
ANDI	I	AND Immediate	$rd = rs1 \& \text{sign_ext}(imm)$	Bitwise AND the sign-extended immediate to register $rs1$ and store result in rd .
JAL	J	Jump		$PC = PC + \text{sign_ext}(imm)$ and store the result in rd .
BEQ	B	Branch		($PC = PC + \text{sign_ext}(imm)$) if $rs1$ is equal to $rs2$.
BNE	B	Branch		($PC = PC + \text{sign_ext}(imm)$) if $rs1$ is not equal to $rs2$.
LW	I	Load		Load word from memory address $[rs1 + \text{sign_ext}(imm)]$ and store it in rd .
SW	S	Store		Store word from register $rs2$ to memory address $[rs1 + \text{sign_ext}(imm)]$.
HALT	-	Halt execution		

Instruction encoding:

Mnemonic	Bit Fields						
	31:27	26:25	24:20	19:15	14:12	11:7	6:0
ADD	0000000		rs2	rs1	000	rd	0110011
SUB	0100000		rs2	rs1	000	rd	0110011
XOR	0000000		rs2	rs1	100	rd	0110011

OR	0000000	rs2	rs1	110	rd	0110011
AND	0000000	rs2	rs1	111	rd	0110011
ADDI	imm[11:0]		rs1	000	rd	0010011
XORI	imm[11:0]		rs1	100	rd	0010011
ORI	imm[11:0]		rs1	110	rd	0010011
ANDI	imm[11:0]		rs1	111	rd	0010011
JAL	imm[20 10:1 11 19:12]				rd	1101111
BEQ	imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011
BNE	imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011
LW	imm[11:0]	rs1	000	rd		0000011
SW	imm					0100011
HALT						1111111

The simulator should handle the following instructions:

- **Instruction Fetch:** Fetch the instruction from memory at the PC value as address.
- **Instruction Decode:** Decode the instruction format in the table above and generates control signals.
- **Execute:** Perform the operation specified by the instruction.
- **Load/ Store:** Perform the load/store operation.
- **Writeback:** Write the result back to the register file. Note that R0 in RISC-V can only contain the zero value.

Each stage must be precise and must not change the state of the processor to the next stage in the next cycle. Each stage should be able to handle the instruction in the following cycle.

The simulator must be able to deal with two types of hazards.

1. **RAW Hazards:** RAW hazards are dealt with using either only forwarding (if possible) or, if not, using stalling + forwarding. Use EX-ID forwarding and MEM-ID forwarding appropriately.
2. **Control Flow Hazards:** The branch conditions are resolved in the ID/RF stage of the pipeline.

The simulator deals with branch instructions as follows:

1. Branches are always assumed to be NOT TAKEN. That is, when a beq is fetched in the IF stage, the PC is speculatively updated as PC+4.
2. Branch conditions are resolved in the ID/RF stage.

3. If the branch is determined to be not taken in the ID/RF stage (as predicted), then the pipeline proceeds without disruptions. If the branch is determined to be taken, then the speculatively fetched instruction is discarded and the nop bit is set for the ID/RR stage for the next cycle. Then the new instruction is fetched in the next cycle using the new branch PC address.

Tasks:

- 1) Draw the schematic for a single stage processor and fill in your code in the to run the simulator. (20 points)
- 2) Draw the schematic for a five stage pipelined processor and fill in your code to run the simulator. The processor should be able to take care of RAW and control hazards by stalling and forwarding. (20 points)
- 3) Measure and report average CPI, Total execution cycles and Instructions per cycle for both these cores by adding performance counters to the simulator. (5 points)
- 4) Compare the results of the two processor implementations and explain why one is faster than the other. (5 points)
- 5) What optimizations can be made to the processor? (1 point)

Your work will be evaluated by the end of the week before the deadline. (50 points - 5 points for each task)

Useful References:

- More details on RISC-V: <https://riscv.org>
- bitset library for C++: <https://en.cppreference.com/w/cpp/bitset>
- g++: <https://gcc.gnu.org/>
- python: <https://www.python.org/downloads/>

